

QA LEADERSHIP ACADEMY

AI Pilot Template

Run a small, time-boxed AI trial with success measures agreed up front, so the decision to scale is evidence-based.

Purpose

This template structures a small, controlled trial of an AI use case so that the decision to adopt, adjust or drop it rests on evidence rather than enthusiasm. It fixes the hypothesis, the scope, the success and guardrail measures, and the controls in advance, and requires a go / no-go decision at the end. It is how you turn "AI seems useful" into a defensible verdict.

When to use it

- For any AI use the Use-Case Matrix rated high-value, once its Risk Assessment is complete.
- When leadership wants proof of value before committing to a tool or approach.
- To convert informal experimentation into a controlled trial with a clear verdict.

Instructions

1. Write a single, testable hypothesis and the measures that would prove or disprove it.
2. Time-box it (a few weeks) and scope it to one team, one use, and a manageable sample.
3. Define success measures AND guardrail measures (what must NOT get worse — e.g. escaped defects, review effort).
4. Confirm the controls from the Risk Assessment and Governance Checklist are in place before you start.
5. At the end, hold a go / no-go: adopt, adjust and re-pilot, or stop — and record why.

Worked example

A pilot for Northstar of a sanctioned, high-value / low-risk use — AI-assisted drafting of test ideas from acceptance criteria — chosen because it delivers value while the riskier uses wait for controls:

Element	Pilot definition
Hypothesis	AI-drafted test ideas from acceptance criteria will speed up test design without reducing the quality of coverage

Element	Pilot definition
Scope	The Web squad only; one AI tool from the approved list; 10 user stories
Duration	3 iterations (~6 weeks) — indicative time-box, not a commitment
Participants	2 QA engineers (one senior, one ramping) + Sofia as reviewer
Success measures	Test-design time per story trends down; reviewers judge idea quality equal or better
Guardrail measures	No rise in escaped defects on piloted stories; review effort not ballooning; no sensitive data entered
Controls in place	No PII/source pasted; human reviews and owns every kept idea; approved tool only
Decision gate	Go / adjust / stop, decided with evidence at the end of iteration 3

What makes this a real pilot rather than dabbling: the guardrail measures. If test-design time falls but escaped defects rise, the pilot has failed even though it "felt" faster. Deciding that up front stops a good demo from being mistaken for a good result — an important discipline at Northstar, where a previous "automate everything" push over-promised and lost trust.

GUARDRAILS, NOT JUST GOALS

Every AI pilot needs a measure of what must not get worse. Speed that quietly costs quality is a failed pilot dressed up as a success.

Blank reusable version

Element	Your definition
Hypothesis	
Scope (team / tool / sample)	
Duration (time-box)	
Participants	
Success measures	
Guardrail measures (must not worsen)	
Controls in place (from Risk Assessment)	
Decision gate (go / adjust / stop)	

Common mistakes

COMMON MISTAKES

Running a pilot with no measures, so the verdict is a vibe; setting success measures but no guardrails, so hidden quality loss passes as a win; scoping it so broadly it cannot be controlled; skipping the Risk Assessment controls "just for the pilot"; and letting a pilot drift on with no decision date, quietly becoming permanent without ever being judged.

How to interpret the results

- Success on the goal AND held guardrails = adopt; success on the goal but a breached guardrail = adjust and re-pilot, not adopt.
- A pilot that "felt great" but moved no measure is a no-go; the feeling is not the evidence.
- Small sample sizes mean treat the result as directional; scale gradually and keep measuring.
- A clean go on a low-risk use also builds the organisational trust you need to pilot the higher-risk ones later.

INSIDE STLC PRO TIP

Pick your first AI pilot to win: a high-value, low-risk use with clear measures. A well-run, evidence-backed pilot earns you the credibility — and the governance goodwill — to take on the harder, higher-value uses next. A sprawling, unmeasured first pilot burns both.