

QA LEADERSHIP ACADEMY

Final Candidate Decision Matrix

Bring every stage together and weight it against the gap you actually have — revealing that there is rarely one right answer.

Purpose

The decision matrix consolidates every stage — screening, technical, behavioural, references — into one weighted comparison, with the weights driven by the gap on your skills matrix and heatmap. Its real value is not that it picks a winner, but that it makes the trade-offs explicit, so the panel argues about weightings (a strategy question) rather than about people (a preference question).

When to use it

- At the final decision, once all candidates have completed the same stages.
- When a panel is split and needs to surface why they disagree.
- Whenever you want a defensible, evidence-based hiring decision on record.

Instructions

1. List the criteria that matter for this hire, drawn from the JD and your capability gaps.
2. Set weights from the team's real need — and be explicit that changing the weights can change the winner. That is a feature, not a flaw.
3. Bring in each candidate's scores from the earlier stages rather than re-judging from memory.
4. Compute weighted totals, then run a second weighting that reflects a different but legitimate strategy, and see whether the answer moves.
5. Discuss the trade-offs the numbers expose; let the panel decide with eyes open. The matrix informs judgement, it does not replace it.
6. Record the decision and the reasoning, in line with your data-retention policy.

Worked example

Northstar's three finalists, scored 1–5 per criterion. Candidate A: deep automation/SDET specialist. Candidate B: strong all-round senior with fintech domain, leadership and mentoring, weaker automation. Candidate C: mid-level generalist, good API/CI, strong learning potential and communication, less senior.

Weighting 1 — "fix the burning platform" (weight the immediate risk: flaky suite, Dan's key-person risk)

Criterion (weight)	Cand A	Cand B	Cand C
Automation depth (x3)	5	2	3
CI/CD & pipeline (x3)	5	2	3
API / contract testing (x2)	3	3	4
Coaching / spreading ownership (x2)	2	5	3
Communication (x1)	2	5	4
Domain (fintech) (x1)	2	5	2
Weighted total (max 60)	46	39	43

Under this weighting, Candidate A wins — the specialist who can most quickly stabilise the suite and pipeline.

Weighting 2 — "build a resilient team" (weight long-term resilience: reduce key-person risk, grow the team, ease the dev/QA rift)

Criterion (weight)	Cand A	Cand B	Cand C
Automation depth (x2)	5	2	3
CI/CD & pipeline (x2)	5	2	3
API / contract testing (x2)	3	3	4
Coaching / spreading ownership (x3)	2	5	3
Communication (x2)	2	5	4
Domain (fintech) (x2)	2	5	2
Weighted total (max 65)	40	48	43

Under this weighting, Candidate B wins — the collaborator and mentor who reduces the team's people risks, even though they would not personally rebuild the suite fastest. And Candidate C sits second in both scenarios: never the top pick, but the lowest-regret choice if the panel cannot agree on strategy.

The point

There is no universal right answer. A is right if the priority is stopping the bleeding now; B is right if the priority is a resilient, well-mentored team that fixes the culture that produced the flaky suite; C is the safe compromise. The matrix does not tell Northstar who to hire — it forces the panel to decide what problem they are hiring to solve, which is the decision that actually matters. Choose the weighting deliberately, then let it choose the candidate.

HR AND LEGAL NOTE

This is a decision-support tool, not legal advice. Always run your process through your own HR/People team and follow local employment law on data handling, right-to-work, records retention, reasonable adjustments and non-discrimination. Ask every candidate for a role the same questions against the same criteria, record the evidence rather than the impression, and keep your notes factual — they may be disclosable.

How to interpret the results

- If the winner flips when you change legitimate weightings, the honest conclusion is "we must decide our strategy first", not "the tool is broken".
- A candidate who is second under every reasonable weighting is often the lowest-risk hire, even if never anyone's first choice.
- Never let the arithmetic overrule a serious, evidenced concern (a behavioural red flag, a failed reference). The matrix ranks, humans decide.
- Beware tuning the weights until your favourite wins — set them from the gap analysis before you compute, and change them only for a stated strategic reason.

Blank reusable version

Criterion (weight)	Cand 1	Cand 2	Cand 3
Weighted total			

Run it twice under two legitimate weightings and compare. If the answer is stable, decide with confidence; if it flips, decide your strategy first.

Common mistakes

COMMON MISTAKES

Reverse-engineering the weights until the preferred candidate wins. Treating the total as a verdict and switching off your judgement. Re-scoring from memory at the final stage instead of carrying evidence forward. Weighting for a generic "great tester" rather than your actual gap. Ignoring that a flipping winner is telling you the panel has not agreed what the role is for. Overruling a red flag because the number was high.

INSIDE STLC PRO TIP

When the matrix keeps changing its mind, stop tuning the numbers and ask the panel one question: "what is the single most important problem this hire must solve in the next six months?" Agree that, and the weights — and the candidate — usually settle themselves.