

QA LEADERSHIP ACADEMY

Quality Health Scorecard

A one-page, RAG-rated read on the overall health of quality across dimensions — and where to invest next.

Purpose

A quality health scorecard gives a single, honest, at-a-glance view of how healthy quality is across the dimensions that matter — not just testing, but the whole system that produces quality: requirements, coverage, automation, environments, defect discipline and outcomes. It exists to answer two questions a leader asks constantly: "Where do we stand?" and "Where should the next unit of effort go?". Unlike a dashboard aimed at executives, the scorecard is a working diagnostic for the QA leader and the engineering leaders they partner with.

When to use it

- In the first weeks of a new QA leadership role, to establish an honest baseline of what you have inherited.
- Quarterly, to track whether quality health is genuinely improving and to re-target investment.
- When deciding between competing improvement bets (more automation vs better environments vs earlier involvement).
- When you need to show progress over time to a sponsor who wants "faster, safer releases".

Instructions

1. Choose the dimensions that describe your quality system end to end — not just test execution.
2. For each, write a plain-language "what good looks like" so the RAG rating is judged against a standard, not a mood.
3. Rate each dimension Red, Amber or Green with a one-line evidence note. Be honest; a scorecard that is all green on day one is not credible.
4. Add a trend arrow (improving / steady / worsening) so the reader sees direction, not just position.
5. Name the single highest-value action for each Red or worsening dimension, with an owner.
6. Re-run it on a fixed cadence with the same dimensions, so the comparison over time is meaningful.

Worked example

An honest baseline scorecard for Northstar as a QA leader would find it on arrival, drawn from the known current state — centralised QA pulled in late, a 5-day manual regression, a large and flaky UI suite, unstable shared staging, inconsistent requirements, and metrics that measure activity rather than outcomes.

Dimension	What good looks like	RAG	Trend	Evidence / note
Requirements quality	Clear, testable acceptance criteria on every story	Red	Steady	Inconsistent across squads; some hand over a design link and a sentence
QA involvement timing	QA engaged from refinement, not after code-complete	Red	Steady	Centralised QA pulled in late, once stories are code-complete
Risk-based coverage	Critical/High risks explicitly covered and tracked	Amber	Improving	Being introduced; not yet consistent across all squads
Automation health	Trusted, fast, low-flake automation owned by the whole team	Red	Worsening	Large UI suite (~1,800 tests), ~6-hour run, ~25% flaky, owned by QA alone, low trust
Regression efficiency	Fast, largely automated regression of critical journeys	Red	Steady	~5-day manual regression before major releases
Environment & test data	Stable environments with a reliable test-data strategy	Red	Steady	Shared staging unstable; frequently "someone else's data"
Defect management discipline	Consistent severity/priority; defects resolved not hoarded	Amber	Steady	Jira workflow exists but severity/priority used loosely
Escaped defects / outcomes	Low and falling customer-affecting escapes	Red	Worsening	Production defects trending up; escaped defects not yet even measured
Metrics & reporting maturity	Outcome metrics with context, not activity counts	Red	Improving	Currently reports test-case and bug counts; catalogue being introduced
Team capability & career path	Skills matrix and a progression ladder	Amber	Steady	Strong exploratory testers; no QA career framework yet

The scorecard immediately tells the investment story. Automation is Red and worsening, but pouring effort straight into more tests would repeat last year's failed "automate everything" push. The upstream Reds — requirements quality and late QA involvement — are cheaper to move and reduce the defects that make regression painful in the first place. So the first bet is earlier involvement and clearer acceptance criteria, the second is stabilising environments, and only then targeted automation of critical journeys. The scorecard turns a vague "improve quality" into a defensible sequence.

Blank reusable version

Dimension	What good looks like	RAG	Trend	Highest-value next action / owner
Requirements quality				
QA involvement timing				
Risk-based coverage				
Automation health				

Dimension	What good looks like	RAG	Trend	Highest-value next action / owner
Regression efficiency				
Environment & test data				
Defect management discipline				
Escaped defects / outcomes				
Metrics & reporting maturity				
Team capability & career path				

Common mistakes

COMMON MISTAKES

Rating everything green to look competent, which destroys the scorecard's credibility the first time reality disagrees; scoring dimensions on mood rather than against a written "what good looks like"; measuring only test execution and ignoring the upstream dimensions (requirements, involvement timing) that cause most defects; changing the dimensions every quarter so no trend is possible; and producing the scorecard without naming the next action, leaving a diagnosis with no prescription.

- Treating the RAG as a grade on the QA team rather than a health check on the whole quality system.
- Investing in the loudest Red (automation) before the cheapest, highest-leverage one (requirements and early involvement).

How to interpret the results

- Look for upstream Reds first. Poor requirements and late involvement cause downstream pain, so fixing them often turns several other dimensions Amber for less effort.
- A Red that is improving needs patience; a Green that is worsening needs attention now — trend can matter more than position.
- Do not chase all Reds at once. The scorecard's job is to sequence a small number of high-leverage bets, not to justify doing everything.
- Beware the tempting Red (a big flaky automation suite) that would repeat a past failure — check that the root cause is addressed before re-investing.
- Run the same dimensions on the same cadence; the value compounds as a trend that shows a sponsor real, sequenced progress rather than a one-off snapshot.

INSIDE STLC PRO TIP

Present your first scorecard with several honest Reds and a clear plan, not a reassuring wall of green. A new leader who says "here is exactly where we are weak, and here is the order I will fix it in" earns far more trust than one who claims everything is fine — and the honest baseline is what lets you prove, quarter on quarter, that you moved the needle.