

QA LEADERSHIP ACADEMY

Release Risk Assessment

A structured read on whether a specific release is safe to ship, and what the residual risk actually is.

Purpose

A release risk assessment answers one question for one release: "If we ship this now, what could go wrong, how likely is it, and what have we done about it?" It differs from the product risk matrix, which shapes testing up front. This is the pre-flight check — completed close to release, reflecting what you actually found, so leaders can make an informed decision rather than a hopeful one.

Its real value shows under pressure. When someone wants to ship tomorrow, this document replaces the argument "QA says no" with evidence: here is the risk, here is the confidence, here is what mitigates it, here is the call.

When to use it

- Immediately before any significant release, as the input to a go/no-go decision.
- When a release is being accelerated and you must state residual risk honestly and fast.
- When a hotfix must go out and there is no time for a full regression pass.
- After a change to a high-blast-radius area (payments, auth, billing, data migrations).

Instructions

1. State what is actually in the release — the changes, not the epic name. Ambiguity here poisons everything downstream.
2. For each area of change, record what was tested, what was not, and why.
3. Assess residual risk per area using probability × impact (see the Product Risk Matrix), reflecting testing already done.
4. Record mitigations that reduce likelihood or impact: feature flags, staged rollout, monitoring, a rehearsed rollback.
5. Set a confidence level and state your recommendation plainly: Go, Go with conditions, or No-Go.
6. Name the decision owner. QA advises on risk; the accountable owner decides.

Worked example

The scenario: it is late afternoon and Tom Fielding wants the Payments checkout change live tomorrow morning to hit a partner commitment. Shared staging was unstable for most of the day, so testing was compressed. The QA lead fills in the assessment rather than simply resisting.

Field	Entry
Release	Payments checkout v2.3 — new capture-to-order flow, touches legacy billing write path
Requested date	Tomorrow 09:00, ahead of a B2B partner commitment
In scope	Card checkout, order creation, billing record write, purchase analytics event
Out of scope (unchanged)	Refunds, subscription renewals, admin reporting
Tested and passing	Happy-path card checkout; order creation; discount application on web and mobile
Tested with concerns	Billing write path — passed twice, but staging reset mid-test, so the sample is thin
Not tested	Behaviour under concurrent load; failure/reconciliation when the provider succeeds but billing write fails
Blocked by	Shared staging instability for ~half the test window; no reliable load environment

Area	Residual risk	Rating	Mitigation available
Capture-to-order mismatch (charged, no order)	Possible × Severe	High	Ship behind a feature flag to 5% of traffic; alert on capture-without-order; rehearsed rollback
Billing write correctness under load	Possible × Major	High	Staged rollout; reconciliation report run at 1 hour and 24 hours; on-call payments engineer briefed
Analytics purchase event	Possible × Minor	Medium	Accept; verify in production within the first hour
Discount / visual issues	Unlikely × Negligible	Low	Accept

Summary	Value
Overall residual risk	High, concentrated in the money path
Confidence in the assessment	Medium — reduced by compressed, staging-limited testing
Recommendation	Go with conditions
Conditions	1) Flagged 5% rollout, not 100%. 2) Capture-without-order alert live before release. 3) Reconciliation report at +1h and +24h. 4) Named on-call owner. 5) Rollback rehearsed and under 15 minutes.
Decision owner	VP Engineering (Priya Nadar), with Head of Product

The outcome is not "no". It is a shippable path that protects the money while honouring the date — and it makes the conditions, not QA, the thing the business signs up to. If leadership declines the conditions, the recommendation reverts to No-Go, and that too is now on the record.

Blank reusable version

Field	Entry
Release	
Requested date	
In scope	
Out of scope	
Tested and passing	
Tested with concerns	
Not tested	
Blocked by	

Area	Residual risk (prob × impact)	Rating	Mitigation available

Summary	Value
Overall residual risk	
Confidence in the assessment	High / Medium / Low
Recommendation	Go / Go with conditions / No-Go
Conditions	
Decision owner	

Common mistakes

COMMON MISTAKES

Reporting a green pass rate while hiding that half the plan was blocked by staging; conflating "we did not test it" with "it works"; giving a binary Go/No-Go with no conditional middle path, which forces a needless fight; owning the decision yourself instead of advising the accountable owner; and omitting confidence, so a thin, rushed assessment reads as authoritative as a thorough one.

- Listing mitigations that do not actually exist yet ("we can roll back" when rollback has never been tested).
- Assessing risk against the plan rather than against what was really executed.

How to interpret the results

- Read residual risk and confidence together. High residual risk with low confidence is the most dangerous combination and rarely a clean Go.
- Concentrated risk (one High area) is often shippable with a targeted mitigation; scattered High risk across many areas usually is not.
- "Go with conditions" is the professional default when a date matters and the risk is containable — it moves the decision from opinion to a checklist.

- If the mitigations are what make it a Go, the mitigations are mandatory, not optional. Track them as release blockers.
- A No-Go should always state what would change it to a Go, so the path forward is obvious.

INSIDE STLC PRO TIP

Never let "we ran out of time to test it" be silently read as "it is fine". The most valuable line in any release assessment is the honest "Not tested" row. Leaders forgive known, stated risk; they do not forgive discovering after an incident that you knew the gap and stayed quiet.